

Text Mining Assignment

Elena Montenegro, Fernando Freire

May 1, 2019

Contents

1	Modules importation and data loading	2
1.1	Data split	4
2	Part I. Construction of an automatic classifier	4
2.1	Pipelines	5
2.1.1	Find additional stopwords	5
2.1.2	Pipelining methods	5
2.2	Main process with prefixed parameters	7
2.3	Main process with grid search parameters	11
3	Part 2: Construction of a clustering of biology documents	17

1 Modules importation and data loading

Script 1.0.1 (python)

```
1 import warnings
2 warnings.filterwarnings('ignore')
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import pandas as pd
6
7 %matplotlib inline
8 from sklearn.feature_extraction.text import CountVectorizer
9 from sklearn.feature_extraction.text import TfidfTransformer
10
11 from sklearn.naive_bayes import MultinomialNB
12 from sklearn.decomposition import TruncatedSVD# SVD = Singular Value Decomposition
13 from sklearn.model_selection import GridSearchCV
14 from sklearn.feature_extraction.text import CountVectorizer
15 from sklearn.feature_extraction.text import TfidfVectorizer
16 from sklearn.preprocessing import StandardScaler, Normalizer, MinMaxScaler, MaxAbsScaler
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.feature_selection import SelectKBest
19 from sklearn.pipeline import Pipeline
20 from sklearn.model_selection import train_test_split
21 from sklearn import metrics
22 from sklearn.svm import SVC, LinearSVC
23 from sklearn.tree import DecisionTreeClassifier
24 from sklearn.neighbors import KNeighborsClassifier
25 from sklearn import tree
26 from sklearn.feature_extraction import stop_words
27
28 random_state=0
```

Script 1.0.2 (python)

```
1 # Data loading
2 NROWS = 50
3 ## Negative dataset
4 df_neg = pd.read_csv('./practica_clase/PRECISION_MEDICINE/negative_training_abstracts.tsv',
5     → sep='\t',
6     header=None, nrows = NROWS)
7
8 df_neg.columns = ['Accession number', 'Title', 'Abstract']
9 df_neg['Label'] = '0' #'neg'
10
11 display(df_neg.head())
12
13 corpus_neg = list(df_neg['Abstract'].values)
14 ### len(corpus_neg) # 4078
15
16 ## Positive
```

```

16 df_pos = pd.read_csv('./practica_clase/PRECISION_MEDICINE/positive_training_abstracts.tsv',
    ↪ sep='\t',
17                        header=None, nrows = NROWS)
18
19 df_pos.columns = ['Accession number', 'Title', 'Abstract']
20 df_pos['Label'] = '1' # 'pos'
21 display(df_pos.head())
22
23 # Add corpus
24 df_corpus = df_neg.append(df_pos)
25 display(df_corpus.head())
26
27 # len(corpus) # 8156
28
29 labels = df_corpus['Label']
30 corpus = df_corpus['Abstract']
31 # len(labels) # 8156
32
33 print(len(corpus), len(labels))

```

	Accession number	Title \
0	29606186	Can reactivity and regulation in infancy predi...
1	29471205	Fabrication of bioinspired, self-cleaning supe...
2	29175165	Functional properties of chickpea protein isol...
3	29098524	Mechanical dyssynchrony alters left ventricula...
4	27507285	Reducing the width of confidence intervals for...

	Abstract	Label
0	A need to identify early infant markers of lat...	0
1	The mechanical properties, corrosion-resistanc...	0
2	In the present study, the effect of Refractanc...	0
3	The impact of left bundle branch block (LBBB) ...	0
4	In the last decade, it has been shown that an ...	0

	Accession number	Title \
0	27829177	A naturally occurring variant of HPV-16 E7 exe...
1	27806271	Functional Analysis of Orail Concatemers Suppo...
2	27796307	KAT2A/KAT2B-targeted acetylome reveals a role ...
3	27795438	The Cellular DNA Helicase ChlR1 Regulates Chro...
4	27794539	Human R1441C LRRK2 regulates the synaptic vesi...

	Abstract	Label
0	Human Papillomavirus E6 and E7 play critical r...	1
1	Store-operated Ca(2+) entry occurs through the...	1
2	Lysine acetylation is a widespread post-transl...	1
3	In papillomavirus infections, the viral genome...	1
4	Mutations in leucine-rich repeat kinase 2 (LRR...	1

	Accession number	Title \
0	29606186	Can reactivity and regulation in infancy predi...
1	29471205	Fabrication of bioinspired, self-cleaning supe...
2	29175165	Functional properties of chickpea protein isol...
3	29098524	Mechanical dyssynchrony alters left ventricula...
4	27507285	Reducing the width of confidence intervals for...

	Abstract	Label
0	A need to identify early infant markers of lat...	0
1	The mechanical properties, corrosion-resistanc...	0
2	In the present study, the effect of Refractanc...	0
3	The impact of left bundle branch block (LBBB) ...	0
4	In the last decade, it has been shown that an ...	0

Output

100 100

1.1 Data split

Script 1.1.1 (python)

```
1 TEST_SIZE = 0.33
2 X_train, X_test, y_train, y_test = train_test_split(
3     corpus, labels, test_size=0.33, random_state=random_state)
```

2 Part I. Construction of an automatic classifier

The following parameters can be adjusted in order to try to maximize the quality of the classifier:

- In function TfidfVectorizer:
 - Parameters that affect the vocabulary quality:
 - * List of stopwords (one of the options is setting it to None)
 - * maxfeatures
 - * max_df, min_df
 - Norm (none, 'l1' or 'l2')
- In Latent Semantic Analysis (LSA):
 - n_components
 - not performing LSA
- Classifier model:
 - You can use strategies included in some of the notebooks we used

- * Logistic Regression,
- * Naïve Bayes,
- * decision trees,
- * SVC
- * or others you learnt from the Machine Learning course (k-nn, neural networks, etc.)

The goal is not to check all possible combinations of these parameters but respond to these questions:

- Which tips can you give about constructing an automatic text classifier? What do you recommend to do? What do you recommend not to do?
- What is the best classifier you have obtained?

Your responses to these questions should be illustrated with tables and/or figures and/or screen captures.

2.1 Pipelines

2.1.1 Find additional stopwords

Script 2.1.1 (python)

```

1 def get_top_n_words(corpus, n=None):
2     """
3     List the top n words in a vocabulary according to occurrence in a text corpus.
4     """
5     vec = CountVectorizer().fit(corpus)
6     bag_of_words = vec.transform(corpus)
7     sum_words = bag_of_words.sum(axis=0)
8     words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
9     words_freq = sorted(words_freq, key = lambda x: x[1], reverse=True)
10
11     return words_freq[:n]
12
13 def improve_stop_words(X_train, n=50):
14     """
15     """
16     common_words = [i[0] for i in get_top_n_words(X_train, n)]
17     eng_and_custom_stopwords = set(list(stop_words.ENGLISH_STOP_WORDS) + common_words)
18     print(len(eng_and_custom_stopwords))
19     return eng_and_custom_stopwords

```

2.1.2 Pipelining methods

Script 2.1.2 (python)

```

1 CLASSIFIERS = ['knn', 'dtree', 'nb', 'lr', 'svc', 'lsvc']
2 REDUCERS = ['svd', None]
3 CV = 4
4
5 def create_text_pipeline(reducer='svd', classifier='nb'):
6     """ Create text vectorization pipeline with optional dimensionality reduction"""

```

```

7  assert reducer in REDUCERS, "ERROR: Reducer %s not supported, only %s" % (reducer,
   ↪ REDUCERS)
8  assert classifier in CLASSIFIERS, "ERROR: Classifier %s not supported, only %s" %
   ↪ (classifier, CLASSIFIERS)
9  num_comp = 100
10 pipeline = [
11     ('vect', TfidfVectorizer()),
12     ('scaler', StandardScaler())
13 ]
14
15 # Reduce dimensions
16 if reducer == 'svd':
17     pipeline.append(('red_svd', TruncatedSVD()))
18     pipeline.append(('norm', MinMaxScaler()))
19 elif reducer == 'kbest':
20     pipeline.append(('red_kbest', SelectKBest(k=num_comp)))
21     pipeline.append(('norm', MinMaxScaler()))
22 elif reducer == 'percentile':
23     pipeline.append(('red_percentile', SelectPercentile(f_classif, percentile=num_comp)))
24     pipeline.append(('norm', MinMaxScaler()))
25 elif reducer == None:
26     #pipeline.append(('normalizer', MaxAbsScaler()))
27     pass
28
29 # Classify
30 if classifier == "nb":
31     pipeline.append(('clf_' + classifier, MultinomialNB()))
32 elif classifier == "lr":
33     pipeline.append(('clf_' + classifier, LogisticRegression()))
34 elif classifier == "svc":
35     pipeline.append(('clf_' + classifier, SVC()))
36 elif classifier == "lsvc":
37     pipeline.append(('clf_' + classifier, LinearSVC()))
38 elif classifier == "dtree":
39     pipeline.append(('clf_' + classifier, DecisionTreeClassifier()))
40 elif classifier == "knn":
41     pipeline.append(('clf_' + classifier, KNeighborsClassifier()))
42     KNeighborsClassifier()
43 elif classifier == None:
44     pass
45
46 #print("Pipeline", pipeline)
47 return Pipeline(pipeline)
48
49 def prediction_metrics(X_train, y_train, X_test, y_test, parameters, reducer="svd",
   ↪ classifier="nb"):
50     """
51     """
52     print("### Reducer: %s Classifier: %s" %(reducer, classifier))
53     pipeline = create_text_pipeline(reducer=reducer, classifier=classifier)
54     # Filter params to only the params related with the pipeline steps
55     filtered_params = {}

```

```

56     for param_key in parameters.keys():
57         if param_key.split('__')[0] in pipeline.named_steps.keys():
58             filtered_params[param_key] = parameters[param_key]
59     pipeline.set_params(**filtered_params)
60     pipeline.fit(X_train, y_train)
61     predicted = pipeline.predict(X_test)
62     print(metrics.classification_report(y_test, predicted))
63     print(metrics.confusion_matrix(y_test, predicted))
64
65 def process_classifications(X_train, y_train, X_test, y_test, parameters,
66                             classifiers=CLASSIFIERS, reducers=REDUCERS):
67     """
68     """
69     for classifier in classifiers:
70         for reducer in reducers:
71             prediction_metrics(X_train, y_train, X_test, y_test, parameters, reducer,
72                               ↪ classifier)
73
74 def prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters_grid,
75 ↪ reducer="svd", classifier="nb", cv=CV):
76     """
77     """
78     print("### Reducer: %s Classifier: %s" %(reducer, classifier))
79     pipeline = create_text_pipeline(reducer=reducer, classifier=classifier)
80     # Filter params to only the params related with the pipeline steps
81     filtered_params = {}
82     for param_key in parameters_grid.keys():
83         if param_key.split('__')[0] in pipeline.named_steps.keys():
84             filtered_params[param_key] = parameters_grid[param_key]
85     grid_model = GridSearchCV(pipeline, filtered_params, cv=cv, iid=False)
86     grid_model.fit(X_train, y_train)
87     for param_name in sorted(filtered_params.keys()):
88         print("\t%s: %r" % (param_name, grid_model.best_params_[param_name]))
89     pipeline.set_params(**grid_model.best_params_)
90     pipeline.fit(X_train, y_train)
91     predicted = pipeline.predict(X_test)
92     print(metrics.classification_report(y_test, predicted))
93     print(metrics.confusion_matrix(y_test, predicted))
94
95 def process_classifications_grid(X_train, y_train, X_test, y_test, parameters, cv=CV,
96                                 classifiers=CLASSIFIERS, reducers=REDUCERS):
97     """
98     """
99     for classifier in classifiers:
100         for reducer in reducers:
101             prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters,
102                                   ↪ reducer, classifier, cv=cv)

```

2.2 Main process with prefixed parameters

Script 2.2.1 (python)

```
1 # First set of parameters
2 param_set_1 = {
3     'vect__norm': None,
4     'vect__smooth_idf': True,
5     'vect__sublinear_tf': True,
6     'vect__max_features': 50,
7     'vect__min_df': 6,
8     'vect__stop_words': 'english',
9     'vect__strip_accents' : 'unicode',
10    'vect__analyzer' : 'word',
11    'vect__token_pattern': r'\w{1,}',
12    'vect__ngram_range' : (1, 2),
13    'scaler__with_mean' : False,
14    'vect__norm': 'l2',
15    'red_svd__n_components': 40,
16    'clf_knn__n_neighbors' : 2
17 }
18
19 # More stop words
20 #eng_and_custom_stopwords = improve_stop_words(X_train, 200)
21 #param_set_1['vect__stop_words'] = eng_and_custom_stopwords
22
23 #process_classifications(X_train, y_train, X_test, y_test, param_set_1, reducers=['svd'],
24 → classifiers=['nb'])
25 process_classifications(X_train, y_train, X_test, y_test, param_set_1)
```

Output

```
### Reducer: svd Classifier: knn
      precision    recall  f1-score   support

         0         0.67        0.93        0.78         15
         1         0.92        0.61        0.73         18

   micro avg         0.76        0.76        0.76         33
   macro avg         0.79        0.77        0.76         33
weighted avg         0.80        0.76        0.75         33

[[14  1]
 [ 7 11]]
### Reducer: None Classifier: knn
      precision    recall  f1-score   support

         0         0.79        1.00        0.88         15
         1         1.00        0.78        0.88         18

   micro avg         0.88        0.88        0.88         33
```


macro avg	0.89	0.89	0.88	33
weighted avg	0.90	0.88	0.88	33

```
[[15 0]
 [ 4 14]]
```

```
### Reducer: svd Classifier: dtree
      precision    recall  f1-score   support

      0         0.92     0.80     0.86         15
      1         0.85     0.94     0.89         18
```

micro avg	0.88	0.88	0.88	33
macro avg	0.89	0.87	0.88	33
weighted avg	0.88	0.88	0.88	33

```
[[12 3]
 [ 1 17]]
```

```
### Reducer: None Classifier: dtree
      precision    recall  f1-score   support

      0         1.00     0.87     0.93         15
      1         0.90     1.00     0.95         18
```

micro avg	0.94	0.94	0.94	33
macro avg	0.95	0.93	0.94	33
weighted avg	0.95	0.94	0.94	33

```
[[13 2]
 [ 0 18]]
```

```
### Reducer: svd Classifier: nb
      precision    recall  f1-score   support

      0         0.82     0.93     0.87         15
      1         0.94     0.83     0.88         18
```

micro avg	0.88	0.88	0.88	33
macro avg	0.88	0.88	0.88	33
weighted avg	0.89	0.88	0.88	33

```
[[14 1]
 [ 3 15]]
```

```
### Reducer: None Classifier: nb
      precision    recall  f1-score   support

      0         0.93     0.87     0.90         15
      1         0.89     0.94     0.92         18
```

micro avg	0.91	0.91	0.91	33
macro avg	0.91	0.91	0.91	33
weighted avg	0.91	0.91	0.91	33

```
[[13 2]
```

```

[ 1 17]]
### Reducer: svd Classifier: lr
      precision    recall  f1-score   support

         0         0.81      0.87      0.84         15
         1         0.88      0.83      0.86         18

    micro avg         0.85      0.85      0.85         33
    macro avg         0.85      0.85      0.85         33
   weighted avg         0.85      0.85      0.85         33

```

```

[[13  2]
 [ 3 15]]
### Reducer: None Classifier: lr
      precision    recall  f1-score   support

         0         0.87      0.87      0.87         15
         1         0.89      0.89      0.89         18

    micro avg         0.88      0.88      0.88         33
    macro avg         0.88      0.88      0.88         33
   weighted avg         0.88      0.88      0.88         33

```

```

[[13  2]
 [ 2 16]]
### Reducer: svd Classifier: svc
      precision    recall  f1-score   support

         0         0.45      1.00      0.62         15
         1         0.00      0.00      0.00         18

    micro avg         0.45      0.45      0.45         33
    macro avg         0.23      0.50      0.31         33
   weighted avg         0.21      0.45      0.28         33

```

```

[[15  0]
 [18  0]]
### Reducer: None Classifier: svc
      precision    recall  f1-score   support

         0         0.87      0.87      0.87         15
         1         0.89      0.89      0.89         18

    micro avg         0.88      0.88      0.88         33
    macro avg         0.88      0.88      0.88         33
   weighted avg         0.88      0.88      0.88         33

```

```

[[13  2]
 [ 2 16]]
### Reducer: svd Classifier: lsvc
      precision    recall  f1-score   support

```

```

      0      0.81      0.87      0.84      15
      1      0.88      0.83      0.86      18

micro avg      0.85      0.85      0.85      33
macro avg      0.85      0.85      0.85      33
weighted avg    0.85      0.85      0.85      33

[[13  2]
 [ 3 15]]
### Reducer: None Classifier: lsvc
      precision      recall  f1-score      support

      0      0.81      0.87      0.84      15
      1      0.88      0.83      0.86      18

micro avg      0.85      0.85      0.85      33
macro avg      0.85      0.85      0.85      33
weighted avg    0.85      0.85      0.85      33

[[13  2]
 [ 3 15]]

```

2.3 Main process with grid search parameters

Script 2.3.1 (python)

```
1
```

Script 2.3.2 (python)

```

1 parameters_grid = {
2     'vect__min_df': [5, 6],
3     #'vect__max_df': [10, 11],
4     'vect__stop_words': (None, 'english', eng_and_custom_stopwords),
5     'vect__max_features': [50],
6     #'vect__smooth_idf': [True, False],
7     'vect__norm': ['l1', 'l2', None]
8     #'red_svd__n_components': (50, 100, 200, None),
9     #'red_svd__n_components': (10, 20, 30, None),
10    #'clf_nb__alpha': (1e-1, 1e-2, 1e-3)
11 }
12
13 parameters_grid = {
14     'vect__norm': ['l1', 'l2', None],
15     'vect__smooth_idf': [True],
16     'vect__sublinear_tf': [True],
17     'vect__max_features': [30, 50],
18     'vect__min_df': [5,6],
19     #'vect__max_df': [7,8],
20     'vect__stop_words': [None, 'english', eng_and_custom_stopwords],

```

```

21     'vect__strip_accents' : ['unicode'],
22     'vect__analyzer' : ['word'],
23     'vect__token_pattern': [r'\w{1,}'],
24     'vect__ngram_range' : [(1, 2)],
25     'scaler__with_mean' : [False],
26     'red_svd__n_components': [2,3],
27     'clf_knn__n_neighbors' : [2, 5]
28 }
29
30 eng_and_custom_stopwords = improve_stop_words(X_train, 200)
31 #prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters_grid, reducer='svd',
32   → classifier="knn", cv=2)
33 process_classifications_grid(X_train, y_train, X_test, y_test, parameters_grid, cv=2)

```

Output

```

### Reducer: svd   Classifier: knn
    clf_knn__n_neighbors: 5
    red_svd__n_components: 2
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 30
    vect__min_df: 6
    vect__ngram_range: (1, 2)
    vect__norm: None
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.88        0.93        0.90         15
         1         0.94        0.89        0.91         18

   micro avg       0.91        0.91        0.91         33
   macro avg       0.91        0.91        0.91         33
weighted avg       0.91        0.91        0.91         33

[[14  1]
 [ 2 16]]
### Reducer: None   Classifier: knn
    clf_knn__n_neighbors: 5
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 50
    vect__min_df: 5
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: 'english'

```

```

vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.88        0.93        0.90         15
         1         0.94        0.89        0.91         18

   micro avg         0.91        0.91        0.91         33
   macro avg         0.91        0.91        0.91         33
weighted avg         0.91        0.91        0.91         33

```

```

[[14  1]
 [ 2 16]]

```

```

### Reducer: svd   Classifier: dtree
      red_svd__n_components: 2
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 30
      vect__min_df: 6
      vect__ngram_range: (1, 2)
      vect__norm: None
      vect__smooth_idf: True
      vect__stop_words: 'english'
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

               0         0.93        0.87        0.90         15
               1         0.89        0.94        0.92         18

   micro avg         0.91        0.91        0.91         33
   macro avg         0.91        0.91        0.91         33
weighted avg         0.91        0.91        0.91         33

```

```

[[13  2]
 [ 1 17]]

```

```

### Reducer: None   Classifier: dtree
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 30
      vect__min_df: 5
      vect__ngram_range: (1, 2)
      vect__norm: 'l1'
      vect__smooth_idf: True
      vect__stop_words: 'english'
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

```

0	0.92	0.80	0.86	15
1	0.85	0.94	0.89	18
micro avg	0.88	0.88	0.88	33
macro avg	0.89	0.87	0.88	33
weighted avg	0.88	0.88	0.88	33

```
[[12 3]
 [ 1 17]]
```

```
### Reducer: svd Classifier: nb
    red_svd__n_components: 2
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 30
    vect__min_df: 5
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

    0         0.82     0.93     0.87         15
    1         0.94     0.83     0.88         18

    micro avg       0.88     0.88     0.88         33
    macro avg       0.88     0.88     0.88         33
    weighted avg     0.89     0.88     0.88         33
```

```
[[14 1]
 [ 3 15]]
```

```
### Reducer: None Classifier: nb
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 50
    vect__min_df: 6
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

    0         0.93     0.87     0.90         15
    1         0.89     0.94     0.92         18

    micro avg       0.91     0.91     0.91         33
    macro avg       0.91     0.91     0.91         33
```

weighted avg	0.91	0.91	0.91	33
--------------	------	------	------	----

```
[[13  2]
 [ 1 17]]
```

```
### Reducer: svd   Classifier: lr
    red_svd__n_components: 2
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 50
    vect__min_df: 6
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

         0         0.93      0.93      0.93        15
         1         0.94      0.94      0.94        18

    micro avg         0.94      0.94      0.94        33
    macro avg         0.94      0.94      0.94        33
    weighted avg         0.94      0.94      0.94        33
```

```
[[14  1]
 [ 1 17]]
```

```
### Reducer: None   Classifier: lr
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 30
    vect__min_df: 5
    vect__ngram_range: (1, 2)
    vect__norm: None
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

         0         0.93      0.87      0.90        15
         1         0.89      0.94      0.92        18

    micro avg         0.91      0.91      0.91        33
    macro avg         0.91      0.91      0.91        33
    weighted avg         0.91      0.91      0.91        33
```

```
[[13  2]
 [ 1 17]]
```

```
### Reducer: svd   Classifier: svc
```

```

red_svd__n_components: 3
scaler__with_mean: False
vect__analyzer: 'word'
vect__max_features: 30
vect__min_df: 5
vect__ngram_range: (1, 2)
vect__norm: 'l2'
vect__smooth_idf: True
vect__stop_words: 'english'
vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.88        0.93        0.90         15
         1         0.94        0.89        0.91         18

    micro avg         0.91        0.91        0.91         33
    macro avg         0.91        0.91        0.91         33
weighted avg         0.91        0.91        0.91         33

```

```

[[14  1]
 [ 2 16]]

```

```

### Reducer: None Classifier: svc
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 30
      vect__min_df: 5
      vect__ngram_range: (1, 2)
      vect__norm: None
      vect__smooth_idf: True
      vect__stop_words: 'english'
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

               0         0.94        1.00        0.97         15
               1         1.00        0.94        0.97         18

    micro avg         0.97        0.97        0.97         33
    macro avg         0.97        0.97        0.97         33
weighted avg         0.97        0.97        0.97         33

```

```

[[15  0]
 [ 1 17]]

```

```

### Reducer: svd Classifier: lsvc
      red_svd__n_components: 3
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 30
      vect__min_df: 6

```



```

vect__ngram_range: (1, 2)
vect__norm: None
vect__smooth_idf: True
vect__stop_words: 'english'
vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.88        0.93        0.90         15
         1         0.94        0.89        0.91         18

   micro avg         0.91        0.91        0.91         33
   macro avg         0.91        0.91        0.91         33
weighted avg         0.91        0.91        0.91         33

[[14  1]
 [ 2 16]]
### Reducer: None Classifier: lsvc
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 30
      vect__min_df: 5
      vect__ngram_range: (1, 2)
      vect__norm: None
      vect__smooth_idf: True
      vect__stop_words: 'english'
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

               0         0.93        0.87        0.90         15
               1         0.89        0.94        0.92         18

   micro avg         0.91        0.91        0.91         33
   macro avg         0.91        0.91        0.91         33
weighted avg         0.91        0.91        0.91         33

[[13  2]
 [ 1 17]]

```

3 Part 2: Construction of a clustering of biology documents

We already know the class information in our dataset (positive and negative) but we will test if an automatic clustering system discovers automatically these classes (“labels”). The objective is to learn strategies that will be very useful when we have to cluster unlabeled documents. Therefore, we “hide” this information (the real class) to the clustering algorithm.

The objective in this section is to check what are the parameters that maximize clustering’s quality. The parameters to be taken into account are:

- In function TfidfVectorizer:
 - Vocabulary (larger or smaller)
 - Norm (none, 'l1' or 'l2')
- In Latent Semantic Analysis (LSA):
 - n_components
 - o not performing LSA
- Normalize the data/not normalize it with “Normalizer” (included in the notebook).

The questions to be responded in this part are:

- Which tips can you give about constructing a text clustering with k-means? What do you recommend to do? What do you recommend not to do?
- What is the best clustering you have obtained? The quality of the cluster is the degree of correspondence between real class and assigned cluster. For example:
 - If there are 2 clusters and cluster 0 contains all examples of positive class and cluster 1 contains all examples of negative class, the clustering is perfect.
 - If there are 2 clusters and cluster 1 contains all examples of positive class and cluster 0 contains all examples of negative class, the clustering is also perfect.
 - If there are 2 clusters and cluster 0 contains 50% of examples of positive class and 50% of examples of negative class, and statistics in cluster 1 are similar, the clustering quality is the worst possible.