

# Text Mining Assignment

Elena Montenegro, Fernando Freire

May 1, 2019

## Contents

|          |  |           |
|----------|--|-----------|
| <b>1</b> | <b>Modules importation and data loading</b>                      | <b>2</b>  |
| 1.1      | Data split . . . . .   | 4         |
| <b>2</b> | <b>Part I. Construction of an automatic classifier</b>           | <b>4</b>  |
| 2.1      | Pipelines . . . . .  | 5         |
| 2.1.1    | Find additional stopwords . . . . .                              | 5         |
| 2.1.2    | Pipelining methods . . . . .                                     | 6         |
| 2.2      | Main process with prefixed parameters . . . . .                  | 8         |
| 2.3      | Main process with grid search parameters . . . . .               | 12        |
| <b>3</b> | <b>Part 2: Construction of a clustering of biology documents</b> | <b>18</b> |
| 3.1      | Pipelining . . . . .   | 21        |

# 1 Modules importation and data loading

## Script 1.0.1 (python)

```
1 import warnings
2 warnings.filterwarnings('ignore')
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import pandas as pd
6 import sys
7 %matplotlib inline
8 from sklearn.feature_extraction.text import CountVectorizer
9 from sklearn.feature_extraction.text import TfidfTransformer
10
11 from sklearn.naive_bayes import MultinomialNB
12 from sklearn.decomposition import TruncatedSVD# SVD = Singular Value Descomposition
13 from sklearn.model_selection import GridSearchCV
14 from sklearn.feature_extraction.text import CountVectorizer
15 from sklearn.feature_extraction.text import TfidfVectorizer
16 from sklearn.preprocessing import StandardScaler, Normalizer, MinMaxScaler, MaxAbsScaler
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.feature_selection import SelectKBest
19 from sklearn.pipeline import Pipeline
20 from sklearn.model_selection import train_test_split
21 from sklearn import metrics
22 from sklearn.svm import SVC, LinearSVC
23 from sklearn.tree import DecisionTreeClassifier
24 from sklearn.neighbors import KNeighborsClassifier
25 from sklearn import tree
26 from sklearn.feature_extraction import stop_words
27 from sklearn.base import TransformerMixin
28 from sklearn.cluster import KMeans
29 from sklearn.metrics import calinski_harabaz_score
30 from sklearn.preprocessing import Normalizer, LabelBinarizer, OneHotEncoder
31
32 random_state=0
```

## Script 1.0.2 (python)

```
1 # Data loading
2 NROWS = sys.maxsize
3 ## Negative dataset
4 df_neg = pd.read_csv('./practica_clase/PRECISION_MEDICINE/negative_training_abstracts.tsv',
5     ↪ sep='\t',
6     header=None, nrows = NROWS)
7
8 df_neg.columns = ['Accession number', 'Title', 'Abstract']
9 df_neg['Label'] = '0' #'neg'
10
11 display(df_neg.head())
12 corpus_neg = list(df_neg['Abstract'].values)
```

```

13 ### len(corpus_neg) # 4078
14
15 ## Positive
16 df_pos = pd.read_csv('./practica_clase/PRECISION_MEDICINE/positive_training_abstracts.tsv',
17   → sep='\t',
18   header=None, nrows = NROWS)
19
20 df_pos.columns = ['Accession number', 'Title', 'Abstract']
21 df_pos['Label'] = '1' # 'pos'
22 display(df_pos.head())
23
24 # Add corpus
25 df_corpus = df_neg.append(df_pos)
26 display(df_corpus.head())
27
28 # len(corpus) # 8156
29
30 labels = df_corpus['Label']
31 corpus = df_corpus['Abstract']
32 # len(labels) # 8156
33
34 print(len(corpus), len(labels))

```

|   | Accession number | Title \   |
|---|------------------|---|
| 0 | 29606186         | Can reactivity and regulation in infancy predi... |
| 1 | 29471205         | Fabrication of bioinspired, self-cleaning supe... |
| 2 | 29175165         | Functional properties of chickpea protein isol... |
| 3 | 29098524         | Mechanical dyssynchrony alters left ventricula... |
| 4 | 27507285         | Reducing the width of confidence intervals for... |

|   | Abstract  | Label |
|---|---|-------|
| 0 | A need to identify early infant markers of lat... | 0     |
| 1 | The mechanical properties, corrosion-resistanc... | 0     |
| 2 | In the present study, the effect of Refractanc... | 0     |
| 3 | The impact of left bundle branch block (LBBB) ... | 0     |
| 4 | In the last decade, it has been shown that an ... | 0     |

|   | Accession number | Title \   |
|---|------------------|---|
| 0 | 27829177         | A naturally occurring variant of HPV-16 E7 exe... |
| 1 | 27806271         | Functional Analysis of Orail Concatemers Suppo... |
| 2 | 27796307         | KAT2A/KAT2B-targeted acetylome reveals a role ... |
| 3 | 27795438         | The Cellular DNA Helicase ChlR1 Regulates Chro... |
| 4 | 27794539         | Human R1441C LRRK2 regulates the synaptic vesi... |

|   | Abstract  | Label |
|---|---|-------|
| 0 | Human Papillomavirus E6 and E7 play critical r... | 1     |
| 1 | Store-operated Ca(2+) entry occurs through the... | 1     |
| 2 | Lysine acetylation is a widespread post-transl... | 1     |

|   |   |   |
|---|---|---|
| 3 | In papillomavirus infections, the viral genome... | 1 |
| 4 | Mutations in leucine-rich repeat kinase 2 (LRR... | 1 |

|   | Accession number | Title \   |
|---|------------------|---|
| 0 | 29606186         | Can reactivity and regulation in infancy predi... |
| 1 | 29471205         | Fabrication of bioinspired, self-cleaning supe... |
| 2 | 29175165         | Functional properties of chickpea protein isol... |
| 3 | 29098524         | Mechanical dyssynchrony alters left ventricula... |
| 4 | 27507285         | Reducing the width of confidence intervals for... |

|   | Abstract  | Label |
|---|---|-------|
| 0 | A need to identify early infant markers of lat... | 0     |
| 1 | The mechanical properties, corrosion-resistanc... | 0     |
| 2 | In the present study, the effect of Refractanc... | 0     |
| 3 | The impact of left bundle branch block (LBBB) ... | 0     |
| 4 | In the last decade, it has been shown that an ... | 0     |

## Output

8156 8156

## 1.1 Data split

### Script 1.1.1 (python)

```

1 TEST_SIZE = 0.33
2 X_train, X_test, y_train, y_test = train_test_split(
3     corpus, labels, test_size=0.33, random_state=random_state)

```

## 2 Part I. Construction of an automatic classifier

The following parameters can be adjusted in order to try to maximize the quality of the classifier:

- In function TfidfVectorizer:
  - Parameters that affect the vocabulary quality:
    - \* List of stopwords (one of the options is setting it to None)
    - \* maxfeatures
    - \* max\_df, min\_df
  - Norm (none, 'l1' or 'l2')
- In Latent Semantic Analysis (LSA):
  - n\_components
  - not performing LSA

- Classifier model:
  - You can use strategies included in some of the notebooks we used
    - \* Logistic Regression,
    - \* Naïve Bayes,
    - \* decision trees,
    - \* SVC
    - \* or others you learnt from the Machine Learning course (k-nn, neural networks, etc.)

The goal is not to check all possible combinations of these parameters but respond to these questions:

- Which tips can you give about constructing an automatic text classifier? What do you recommend to do? What do you recommend not to do?
- What is the best classifier you have obtained?

Your responses to these questions should be illustrated with tables and/or figures and/or screen captures.

## 2.1 Pipelines

### 2.1.1 Find additional stopwords

#### Script 2.1.1 (python)

```

1 def get_top_n_words(corpus, n=None):
2     """
3     List the top n words in a vocabulary according to occurrence in a text corpus.
4     """
5     vec = CountVectorizer().fit(corpus)
6     bag_of_words = vec.transform(corpus)
7     sum_words = bag_of_words.sum(axis=0)
8     words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
9     words_freq = sorted(words_freq, key = lambda x: x[1], reverse=True)
10
11     return words_freq[:n]
12
13 def improve_stop_words(X_train, n=50):
14     """
15     """
16     common_words = [i[0] for i in get_top_n_words(X_train, n)]
17     eng_and_custom_stopwords = set(list(stop_words.ENGLISH_STOP_WORDS) + common_words)
18     print(len(eng_and_custom_stopwords))
19     return eng_and_custom_stopwords

```

## 2.1.2 Pipelining methods

### Script 2.1.2 (python)

```
1 CLASSIFIERS = ['knn', 'dtree', 'nb', 'lr', 'svc', 'lsvc']
2 CLASSIFIERS_FROM_CLUSTERS = ['kmeans']
3 REDUCERS = ['svd', None]
4 CV = 4
5
6 class KMeans_foo(KMeans):
7     def fit_transform(self, X, y=None):
8         return self.fit_predict(X)
9
10 class Transformer(TransformerMixin):
11     def __init__(self, model):
12         self.model = model
13
14     def fit(self, y=None, *args, **kwargs):
15         self.model.fit(*args, **kwargs)
16         return self
17
18     def transform(self, X, **transform_params):
19         return pd.DataFrame(self.model.predict(X))
20
21 def create_text_pipeline(reducer='svd', classifier="nb"):
22     """ Create text vectorization pipeline with optional dimensionality reduction"""
23     assert reducer in REDUCERS, "ERROR: Reducer %s not supported, only %s" % (reducer,
24     ↪ REDUCERS)
25     assert classifier in CLASSIFIERS + CLASSIFIERS_FROM_CLUSTERS, \
26         "ERROR: Classifier %s not supported, only %s" % (classifier, CLASSIFIERS +
27     ↪ CLASSIFIERS_FROM_CLUSTERS)
28     pipeline = [
29         ('vect', TfidfVectorizer()),
30         ('scaler', StandardScaler())
31     ]
32     num_comp = 3
33     # Reduce dimensions
34     if reducer == 'svd':
35         pipeline.append(('red_svd', TruncatedSVD()))
36         pipeline.append(('norm', MinMaxScaler()))
37     elif reducer == 'kbest':
38         pipeline.append(('red_kbest', SelectKBest(k=num_comp)))
39         pipeline.append(('norm', MinMaxScaler()))
40     elif reducer == 'percentile':
41         pipeline.append(('red_percentile', SelectPercentile(f_classif, percentile=num_comp)))
42         pipeline.append(('norm', MinMaxScaler()))
43     elif reducer == None:
44         #pipeline.append(('normalizer', MaxAbsScaler()))
45         pass
46
47     # Classify
48     if classifier == "nb":
49         pipeline.append(('clf_' + classifier, MultinomialNB()))
```

```

48 elif classifier == "lr":
49     pipeline.append(('clf_' + classifier, LogisticRegression()))
50 elif classifier == "svc":
51     pipeline.append(('clf_' + classifier, SVC()))
52 elif classifier == "lsvc":
53     pipeline.append(('clf_' + classifier, LinearSVC()))
54 elif classifier == "dtree":
55     pipeline.append(('clf_' + classifier, DecisionTreeClassifier()))
56 elif classifier == "knn":
57     pipeline.append(('clf_' + classifier, KNeighborsClassifier()))
58 elif classifier == "kmeans":
59     pipeline.append(('normalizer', Normalizer()))
60     #pipeline.append(('cluster_kmeans' + classifier, Transformer(KMeans_foo(2))))
61     pipeline.append(('cluster_kmeans', KMeans(2)))
62     pipeline.append(('one_hot_encoder', OneHotEncoder(sparse=False)))
63     #pipeline.append(('binarize', LabelBinarizer()))
64 elif classifier == None:
65     pass
66
67 print("Pipeline", pipeline)
68 return Pipeline(pipeline)
69
70 def prediction_metrics(X_train, y_train, X_test, y_test, parameters, reducer="svd",
71 ↪ classifier="nb"):
72     """
73     """
74     print("### Reducer: %s Classifier: %s" %(reducer, classifier))
75     pipeline = create_text_pipeline(reducer=reducer, classifier=classifier)
76     # Filter params to only the params related with the pipeline steps
77     filtered_params = {}
78     for param_key in parameters.keys():
79         if param_key.split('__')[0] in pipeline.named_steps.keys():
80             filtered_params[param_key] = parameters[param_key]
81     pipeline.set_params(**filtered_params)
82     pipeline.fit(X_train, y_train)
83     if classifier == "kmeans":
84         predicted = pipeline.fit_transform(X_test)
85         print(predicted)
86         print(y_test)
87     else:
88         predicted = pipeline.predict(X_test)
89         print(metrics.classification_report(y_test, predicted))
90         print(metrics.confusion_matrix(y_test, predicted))
91
92 def process_classifications(X_train, y_train, X_test, y_test, parameters,
93 ↪ classifiers=CLASSIFIERS, reducers=REDUCERS):
94     """
95     """
96     for classifier in classifiers:
97         for reducer in reducers:
98             prediction_metrics(X_train, y_train, X_test, y_test, parameters, reducer,
99 ↪ classifier)

```

```

98
99 def prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters_grid,
   ⇨ reducer="svd", classifier="nb", cv=CV):
100     """
101     """
102     print("### Reducer: %s Classifier: %s" %(reducer, classifier))
103     pipeline = create_text_pipeline(reducer=reducer, classifier=classifier)
104     # Filter params to only the params related with the pipeline steps
105     filtered_params = {}
106     for param_key in parameters_grid.keys():
107         if param_key.split('__')[0] in pipeline.named_steps.keys():
108             filtered_params[param_key] = parameters_grid[param_key]
109     grid_model = GridSearchCV(pipeline, filtered_params, cv=cv, iid=False)
110     grid_model.fit(X_train, y_train)
111     for param_name in sorted(filtered_params.keys()):
112         print("\t%s: %r" % (param_name, grid_model.best_params_[param_name]))
113     pipeline.set_params(**grid_model.best_params_)
114     pipeline.fit(X_train, y_train)
115     predicted = pipeline.predict(X_test)
116     print(metrics.classification_report(y_test, predicted))
117     print(metrics.confusion_matrix(y_test, predicted))
118
119 def process_classifications_grid(X_train, y_train, X_test, y_test, parameters, cv=CV,
120                                classifiers=CLASSIFIERS, reducers=REDUCERS):
121     """
122     """
123     for classifier in classifiers:
124         for reducer in reducers:
125             prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters,
   ⇨ reducer, classifier, cv=cv)

```

## 2.2 Main process with prefixed parameters

### Script 2.2.1 (python)

```

1  # First set of parameters
2  param_set_1 = {
3      'vect__norm': None,
4      'vect__smooth_idf': True,
5      'vect__sublinear_tf': True,
6      'vect__max_features': 1000,
7      'vect__min_df': 6,
8      'vect__stop_words': 'english',
9      'vect__strip_accents': 'unicode',
10     'vect__analyzer': 'word',
11     'vect__token_pattern': r'\w{1,}',
12     'vect__ngram_range': (1, 2),
13     'scaler__with_mean': False,
14     'vect__norm': 'l2',
15     'red_svd__n_components': 40,

```



```

16     'clf_knn__n_neighbors' : 2
17 }
18
19 # More stop words
20 #eng_and_custom_stopwords = improve_stop_words(X_train, 200)
21 #param_set_1['vect__stop_words'] = eng_and_custom_stopwords
22
23 #process_classifications(X_train, y_train, X_test, y_test, param_set_1, reducers=['svd'],
24   → classifiers=['nb'])
25 process_classifications(X_train, y_train, X_test, y_test, param_set_1)

```

## Output

```

### Reducer: svd   Classifier: knn
      precision    recall  f1-score   support

         0         0.68      1.00      0.81         15
         1         1.00      0.61      0.76         18

   micro avg       0.79      0.79      0.79         33
   macro avg       0.84      0.81      0.78         33
weighted avg       0.86      0.79      0.78         33

```

```

[[15  0]
 [ 7 11]]
### Reducer: None   Classifier: knn
      precision    recall  f1-score   support

         0         0.60      1.00      0.75         15
         1         1.00      0.44      0.62         18

   micro avg       0.70      0.70      0.70         33
   macro avg       0.80      0.72      0.68         33
weighted avg       0.82      0.70      0.68         33

```

```

[[15  0]
 [10  8]]
### Reducer: svd   Classifier: dtree
      precision    recall  f1-score   support

         0         0.83      1.00      0.91         15
         1         1.00      0.83      0.91         18

   micro avg       0.91      0.91      0.91         33
   macro avg       0.92      0.92      0.91         33
weighted avg       0.92      0.91      0.91         33

```

```

[[15  0]
 [ 3 15]]
### Reducer: None   Classifier: dtree

```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.87   | 0.93     | 15      |
| 1            | 0.90      | 1.00   | 0.95     | 18      |
| micro avg    | 0.94      | 0.94   | 0.94     | 33      |
| macro avg    | 0.95      | 0.93   | 0.94     | 33      |
| weighted avg | 0.95      | 0.94   | 0.94     | 33      |

[[13 2]

[ 0 18]]

### Reducer: svd Classifier: nb

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.83      | 1.00   | 0.91     | 15      |
| 1            | 1.00      | 0.83   | 0.91     | 18      |
| micro avg    | 0.91      | 0.91   | 0.91     | 33      |
| macro avg    | 0.92      | 0.92   | 0.91     | 33      |
| weighted avg | 0.92      | 0.91   | 0.91     | 33      |

[[15 0]

[ 3 15]]

### Reducer: None Classifier: nb

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.93      | 0.87   | 0.90     | 15      |
| 1            | 0.89      | 0.94   | 0.92     | 18      |
| micro avg    | 0.91      | 0.91   | 0.91     | 33      |
| macro avg    | 0.91      | 0.91   | 0.91     | 33      |
| weighted avg | 0.91      | 0.91   | 0.91     | 33      |

[[13 2]

[ 1 17]]

### Reducer: svd Classifier: lr

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 1.00   | 0.97     | 15      |
| 1            | 1.00      | 0.94   | 0.97     | 18      |
| micro avg    | 0.97      | 0.97   | 0.97     | 33      |
| macro avg    | 0.97      | 0.97   | 0.97     | 33      |
| weighted avg | 0.97      | 0.97   | 0.97     | 33      |

[[15 0]

[ 1 17]]

### Reducer: None Classifier: lr

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.93      | 0.87   | 0.90     | 15      |
| 1 | 0.89      | 0.94   | 0.92     | 18      |

|              |      |      |      |    |
|--------------|------|------|------|----|
| micro avg    | 0.91 | 0.91 | 0.91 | 33 |
| macro avg    | 0.91 | 0.91 | 0.91 | 33 |
| weighted avg | 0.91 | 0.91 | 0.91 | 33 |

```
[[13 2]
 [ 1 17]]
```

```
### Reducer: svd Classifier: svc
precision recall f1-score support
```

|   |      |      |      |    |
|---|------|------|------|----|
| 0 | 0.45 | 1.00 | 0.62 | 15 |
| 1 | 0.00 | 0.00 | 0.00 | 18 |

|              |      |      |      |    |
|--------------|------|------|------|----|
| micro avg    | 0.45 | 0.45 | 0.45 | 33 |
| macro avg    | 0.23 | 0.50 | 0.31 | 33 |
| weighted avg | 0.21 | 0.45 | 0.28 | 33 |

```
[[15 0]
 [18 0]]
```

```
### Reducer: None Classifier: svc
precision recall f1-score support
```

|   |      |      |      |    |
|---|------|------|------|----|
| 0 | 0.93 | 0.87 | 0.90 | 15 |
| 1 | 0.89 | 0.94 | 0.92 | 18 |

|              |      |      |      |    |
|--------------|------|------|------|----|
| micro avg    | 0.91 | 0.91 | 0.91 | 33 |
| macro avg    | 0.91 | 0.91 | 0.91 | 33 |
| weighted avg | 0.91 | 0.91 | 0.91 | 33 |

```
[[13 2]
 [ 1 17]]
```

```
### Reducer: svd Classifier: lsvc
precision recall f1-score support
```

|   |      |      |      |    |
|---|------|------|------|----|
| 0 | 0.93 | 0.93 | 0.93 | 15 |
| 1 | 0.94 | 0.94 | 0.94 | 18 |

|              |      |      |      |    |
|--------------|------|------|------|----|
| micro avg    | 0.94 | 0.94 | 0.94 | 33 |
| macro avg    | 0.94 | 0.94 | 0.94 | 33 |
| weighted avg | 0.94 | 0.94 | 0.94 | 33 |

```
[[14 1]
 [ 1 17]]
```

```
### Reducer: None Classifier: lsvc
precision recall f1-score support
```

|   |      |      |      |    |
|---|------|------|------|----|
| 0 | 0.93 | 0.87 | 0.90 | 15 |
| 1 | 0.89 | 0.94 | 0.92 | 18 |

|              |      |      |      |    |
|--------------|------|------|------|----|
| micro avg    | 0.91 | 0.91 | 0.91 | 33 |
| macro avg    | 0.91 | 0.91 | 0.91 | 33 |
| weighted avg | 0.91 | 0.91 | 0.91 | 33 |

```
[[13  2]
 [ 1 17]]
```

## 2.3 Main process with grid search parameters

Script 2.3.1 (python)

```
1 parameters_grid = {
2     'vect__min_df': [5, 6],
3     #'vect__max_df': [10, 11],
4     'vect__stop_words': (None, 'english', eng_and_custom_stopwords),
5     'vect__max_features': [50],
6     #'vect__smooth_idf': [True, False],
7     'vect__norm': ['l1', 'l2', None]
8     #'red_svd__n_components': (50, 100, 200, None),
9     #'red_svd__n_components': (10, 20, 30, None),
10    #'clf_nb__alpha': (1e-1, 1e-2, 1e-3)
11 }
12
13 parameters_grid = {
14     'vect__norm': ['l1', 'l2', None],
15     'vect__smooth_idf': [True],
16     'vect__sublinear_tf': [True],
17     'vect__max_features': [900, 1000],
18     'vect__min_df': [5, 6],
19     #'vect__max_df': [7, 8],
20     'vect__stop_words': [None, 'english', eng_and_custom_stopwords],
21     'vect__strip_accents' : ['unicode'],
22     'vect__analyzer' : ['word'],
23     'vect__token_pattern': [r'\w{1,}'],
24     'vect__ngram_range' : [(1, 2)],
25     'scaler__with_mean' : [False],
26     'red_svd__n_components': [2, 30, 40],
27     'clf_knn__n_neighbors' : [2, 5]
28 }
29
30 eng_and_custom_stopwords = improve_stop_words(X_train, 200)
31 #prediction_metrics_grid(X_train, y_train, X_test, y_test, parameters_grid, reducer='svd',
32     → classifier="knn", cv=2)
33 process_classifications_grid(X_train, y_train, X_test, y_test, parameters_grid, cv=2)
```

### Output

```
448
### Reducer: svd   Classifier: knn
    clf_knn__n_neighbors: 5
    red_svd__n_components: 2
    scaler__with_mean: False
    vect__analyzer: 'word'
```

```

vect__max_features: 900
vect__min_df: 5
vect__ngram_range: (1, 2)
vect__norm: 'l2'
vect__smooth_idf: True
vect__stop_words: None
vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.96      0.91      0.93        163
         1         0.92      0.96      0.94        167

 micro avg       0.94      0.94      0.94       330
 macro avg       0.94      0.94      0.94       330
weighted avg       0.94      0.94      0.94       330

```

```

[[149 14]
 [ 7 160]]

```

```

### Reducer: None Classifier: knn
      clf_knn__n_neighbors: 5
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 1000
      vect__min_df: 6
      vect__ngram_range: (1, 2)
      vect__norm: 'l2'
      vect__smooth_idf: True
      vect__stop_words: None
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

               0         0.94      0.93      0.94        163
               1         0.93      0.95      0.94        167

 micro avg       0.94      0.94      0.94       330
 macro avg       0.94      0.94      0.94       330
weighted avg       0.94      0.94      0.94       330

```

```

[[152 11]
 [ 9 158]]

```

```

### Reducer: svd Classifier: dtree
      red_svd__n_components: 30
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 1000
      vect__min_df: 6
      vect__ngram_range: (1, 2)
      vect__norm: 'l1'

```

```

vect__smooth_idf: True
vect__stop_words: 'english'
vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.94      0.92      0.93        163
         1         0.92      0.94      0.93        167

 micro avg       0.93      0.93      0.93       330
 macro avg       0.93      0.93      0.93       330
weighted avg       0.93      0.93      0.93       330

```

```

[[150 13]
 [ 10 157]]

```

```

### Reducer: None Classifier: dtree
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 900
      vect__min_df: 6
      vect__ngram_range: (1, 2)
      vect__norm: None
      vect__smooth_idf: True
      vect__stop_words: None
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'
            precision    recall  f1-score   support

               0         0.83      0.84      0.83        163
               1         0.84      0.83      0.83        167

 micro avg       0.83      0.83      0.83       330
 macro avg       0.83      0.83      0.83       330
weighted avg       0.83      0.83      0.83       330

```

```

[[137 26]
 [ 29 138]]

```

```

### Reducer: svd Classifier: nb
      red_svd__n_components: 30
      scaler__with_mean: False
      vect__analyzer: 'word'
      vect__max_features: 1000
      vect__min_df: 5
      vect__ngram_range: (1, 2)
      vect__norm: 'l2'
      vect__smooth_idf: True
      vect__stop_words: None
      vect__strip_accents: 'unicode'
      vect__sublinear_tf: True
      vect__token_pattern: '\\w{1,}'

```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.95      | 0.96   | 0.96     | 163     |
| 1            | 0.96      | 0.95   | 0.96     | 167     |
| micro avg    | 0.96      | 0.96   | 0.96     | 330     |
| macro avg    | 0.96      | 0.96   | 0.96     | 330     |
| weighted avg | 0.96      | 0.96   | 0.96     | 330     |

[[157 6]

[ 8 159]]

### Reducer: None Classifier: nb

scaler\_\_with\_mean: False

vect\_\_analyzer: 'word'

vect\_\_max\_features: 1000

vect\_\_min\_df: 6

vect\_\_ngram\_range: (1, 2)

vect\_\_norm: 'l2'

vect\_\_smooth\_idf: True

vect\_\_stop\_words: 'english'

vect\_\_strip\_accents: 'unicode'

vect\_\_sublinear\_tf: True

vect\_\_token\_pattern: '\\w{1,}'

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.88   | 0.92     | 163     |
| 1            | 0.89      | 0.97   | 0.93     | 167     |
| micro avg    | 0.92      | 0.92   | 0.92     | 330     |
| macro avg    | 0.93      | 0.92   | 0.92     | 330     |
| weighted avg | 0.93      | 0.92   | 0.92     | 330     |

[[143 20]

[ 5 162]]

### Reducer: svd Classifier: lr

red\_svd\_\_n\_components: 40

scaler\_\_with\_mean: False

vect\_\_analyzer: 'word'

vect\_\_max\_features: 900

vect\_\_min\_df: 5

vect\_\_ngram\_range: (1, 2)

vect\_\_norm: 'l2'

vect\_\_smooth\_idf: True

vect\_\_stop\_words: None

vect\_\_strip\_accents: 'unicode'

vect\_\_sublinear\_tf: True

vect\_\_token\_pattern: '\\w{1,}'

|  | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

|   |      |      |      |     |
|---|------|------|------|-----|
| 0 | 0.97 | 0.95 | 0.96 | 163 |
|---|------|------|------|-----|

|   |      |      |      |     |
|---|------|------|------|-----|
| 1 | 0.95 | 0.97 | 0.96 | 167 |
|---|------|------|------|-----|

|              |      |      |      |     |
|--------------|------|------|------|-----|
| micro avg    | 0.96 | 0.96 | 0.96 | 330 |
| macro avg    | 0.96 | 0.96 | 0.96 | 330 |
| weighted avg | 0.96 | 0.96 | 0.96 | 330 |

```
[[155  8]
 [ 5 162]]
```

```
### Reducer: None Classifier: lr
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 1000
    vect__min_df: 6
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: None
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

            0         0.98        0.94        0.96         163
            1         0.94        0.98        0.96         167

    micro avg         0.96        0.96        0.96         330
    macro avg         0.96        0.96        0.96         330
    weighted avg         0.96        0.96        0.96         330
```

```
[[153 10]
 [ 3 164]]
```

```
### Reducer: svd Classifier: svc
    red_svd__n_components: 40
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 900
    vect__min_df: 5
    vect__ngram_range: (1, 2)
    vect__norm: 'l2'
    vect__smooth_idf: True
    vect__stop_words: None
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
        precision    recall  f1-score   support

            0         0.97        0.94        0.96         163
            1         0.95        0.98        0.96         167

    micro avg         0.96        0.96        0.96         330
    macro avg         0.96        0.96        0.96         330
    weighted avg         0.96        0.96        0.96         330
```

```
[[154  9]
```



```

[ 4 163]]
### Reducer: None Classifier: svc
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 900
    vect__min_df: 6
    vect__ngram_range: (1, 2)
    vect__norm: None
    vect__smooth_idf: True
    vect__stop_words: 'english'
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.97    0.94    0.96         163
         1         0.95    0.98    0.96         167

    micro avg         0.96    0.96    0.96        330
    macro avg         0.96    0.96    0.96        330
weighted avg         0.96    0.96    0.96        330

```

```

[[154 9]
[ 4 163]]
### Reducer: svd Classifier: lsvc
    red_svd__n_components: 40
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 900
    vect__min_df: 5
    vect__ngram_range: (1, 2)
    vect__norm: 'l1'
    vect__smooth_idf: True
    vect__stop_words: None
    vect__strip_accents: 'unicode'
    vect__sublinear_tf: True
    vect__token_pattern: '\\w{1,}'
      precision    recall  f1-score   support

         0         0.96    0.96    0.96         163
         1         0.96    0.96    0.96         167

    micro avg         0.96    0.96    0.96        330
    macro avg         0.96    0.96    0.96        330
weighted avg         0.96    0.96    0.96        330

```

```

[[156 7]
[ 6 161]]
### Reducer: None Classifier: lsvc
    scaler__with_mean: False
    vect__analyzer: 'word'
    vect__max_features: 1000

```

```

vect__min_df: 6
vect__ngram_range: (1, 2)
vect__norm: 'l2'
vect__smooth_idf: True
vect__stop_words: None
vect__strip_accents: 'unicode'
vect__sublinear_tf: True
vect__token_pattern: '\\w{1,}'

```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.98      | 0.94   | 0.96     | 163     |
| 1            | 0.94      | 0.98   | 0.96     | 167     |
| micro avg    | 0.96      | 0.96   | 0.96     | 330     |
| macro avg    | 0.96      | 0.96   | 0.96     | 330     |
| weighted avg | 0.96      | 0.96   | 0.96     | 330     |

```

[[153 10]
 [ 3 164]]

```

### 3 Part 2: Construction of a clustering of biology documents

We already know the class information in our dataset (positive and negative) but we will test if an automatic clustering system discovers automatically these classes (“labels”). The objective is to learn strategies that will be very useful when we have to cluster unlabeled documents. Therefore, we “hide” this information (the real class) to the clustering algorithm.

The objective in this section is to check what are the parameters that maximize clustering’s quality. The parameters to be taken into account are:

- In function TfidfVectorizer:
  - Vocabulary (larger or smaller)
  - Norm (none, ‘l1’ or ‘l2’)
- In Latent Semantic Analysis (LSA):
  - n\_components
  - o not performing LSA
- Normalize the data/not normalize it with “Normalizer” (included in the notebook).

The questions to be responded in this part are:

- Which tips can you give about constructing a text clustering with k-means? What do you recommend to do? What do you recommend not to do?
- What is the best clustering you have obtained? The quality of the cluster is the degree of correspondence between real class and assigned cluster. For example:

- If there are 2 clusters and cluster 0 contains all examples of positive class and cluster 1 contains all examples of negative class, the clustering is perfect.
- If there are 2 clusters and cluster 1 contains all examples of positive class and cluster 0 contains all examples of negative class, the clustering is also perfect.
- If there are 2 clusters and cluster 0 contains 50% of examples of positive class and 50% of examples of negative class, and statistics in cluster 1 are similar, the clustering quality is the worst possible.

### Script 3.0.1 (python)

```

1 from sklearn.cluster import KMeans
2 from sklearn.metrics import calinski_harabaz_score
3 from sklearn.preprocessing import Normalizer
4 from sklearn.pipeline import make_pipeline
5 from sklearn.preprocessing import Normalizer
6
7 vectorizador = TfidfVectorizer(max_df=1., max_features=1000, norm='l2',
8                               min_df=1, stop_words='english',
9                               #stop_words=stopwords,
10                              #token_pattern=r'(?u)\b[A-Za-z]+\b',
11                              #token_pattern=r'(?ui)\b\w*[a-z]+\w*\b',
12                              use_idf=True)
13 X = vectorizador.fit_transform(X_test)
14
15 print(X.shape)
16 n_componentes = 100
17 svd_truncado = TruncatedSVD(n_componentes)
18 normalizador = Normalizer(copy=False)
19
20 lsa = make_pipeline(svd_truncado, normalizador)
21 #lsa = svd_truncado
22
23 X_lsa = lsa.fit_transform(X)
24
25 varianza_explicada = svd_truncado.explained_variance_ratio_.sum()
26 normalizer = Normalizer()
27 X_lsa_norm = normalizer.fit_transform(X_lsa)
28 X_km = X_lsa_norm
29
30 qmetric = calinski_harabaz_score
31
32 Nclusters_max = 15
33 Nrepetitions = 100
34
35 qualities = []
36 inertias = []
37 models = []
38 kini = 1
39 kfin = 4
40 for k in range(kini, kfin+1):
41     print("Evaluando k=%d" % k)
42     km = KMeans(n_clusters=k,
43                 init='k-means++', n_init=Nrepetitions,

```

```

44         max_iter=500, random_state=2)
45     km.fit(X_km)
46     models.append(km)
47     inertias.append(km.inertia_)
48     if k > 1:
49         qualities.append(qmetric(X_km, km.labels_))
50     else:
51         qualities.append(0)

```

## Output

```

(2692, 1000)
Evaluando k=1
Evaluando k=2
Evaluando k=3
Evaluando k=4



```

## Script 3.0.2 (python)

```

1  fig = plt.figure(figsize=(14,3))
2
3  ax = plt.subplot(1,2,1)
4  plt.plot(range(kini,kfin+1), inertias, marker='o')
5  plt.xlabel('number of clusters')
6  plt.title('clustering inertia')
7
8  ax = plt.subplot(1,2,2)
9  plt.plot(range(kini,kfin+1), qualities, marker='o')
10 plt.xlabel('number of clusters')
11 plt.title('clustering quality')
12 plt.show()
13
14 best = pd.Series(qualities).idxmax() # get index for the best model
15 print("Best number of clusters", best)
16 km = models[best]
17 n_clusters = km.get_params()['n_clusters']
18 clusters = km.labels_
19 print ('Number of clusters of best quality', n_clusters)

```



## Output

Best number of clusters 1  
Number of clusters of best quality 2

## Script 3.0.3 (python)

```
1 # We choose the best option to evaluate the quality of prediction
2
3 # First we try with labels as is
4 labels_predicted = [str(label) for label in km.labels_]
5 predicted = pd.Series(labels_predicted)
6 print(metrics.classification_report(y_test, predicted))
7 print(metrics.confusion_matrix(y_test, predicted))
8
9 # Alternatively we invert the label to match the real labels of each group
10 labels_predicted = [str((label + 1)%2) for label in km.labels_]
11 predicted = pd.Series(labels_predicted)
12 print(metrics.classification_report(y_test, predicted))
13 print(metrics.confusion_matrix(y_test, predicted))
```

## Output

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.99      | 0.87   | 0.93     | 1348    |
| 1            | 0.88      | 0.99   | 0.93     | 1344    |
| micro avg    | 0.93      | 0.93   | 0.93     | 2692    |
| macro avg    | 0.94      | 0.93   | 0.93     | 2692    |
| weighted avg | 0.94      | 0.93   | 0.93     | 2692    |

[[1169 179]  
[ 7 1337]]

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.12      | 0.13   | 0.12     | 1348    |
| 1            | 0.01      | 0.01   | 0.01     | 1344    |
| micro avg    | 0.07      | 0.07   | 0.07     | 2692    |
| macro avg    | 0.06      | 0.07   | 0.07     | 2692    |
| weighted avg | 0.06      | 0.07   | 0.07     | 2692    |

[[ 179 1169]  
[1337 7]]

## 3.1 Pipelining

Can I put all in the pipeline defined previously?

### Script 3.1.1 (python)

```
1 # First set of parameters
2 param_set_1 = {
3     'vect__norm': None,
4     'vect__smooth_idf': True,
5     'vect__sublinear_tf': True,
6     'vect__max_features': 1000,
7     'vect__min_df': 6,
8     'vect__stop_words': 'english',
9     'vect__strip_accents' : 'unicode',
10    'vect__analyzer' : 'word',
11    'vect__token_pattern': r'\w{1,}',
12    'vect__ngram_range' : (1, 2),
13    'scaler__with_mean' : False,
14    'vect__norm': 'l2',
15    'red_svd__n_components': 40,
16    'clf_knn__n_neighbors' : 2
17 }
18
19 # More stop words
20 #eng_and_custom_stopwords = improve_stop_words(X_train, 200)
21 #param_set_1['vect__stop_words'] = eng_and_custom_stopwords
22
23 process_classifications(X_train, y_train, X_test, y_test, param_set_1, reducers=['svd'],
24     ↪ classifiers=['kmeans'])
25
26 #process_classifications(X_train, y_train, X_test, y_test, param_set_1)
```